

Assignment3_45731365_StanleyEpuna_COMP8210

October 29, 2021

0.0.1 Importing / installing libraries

```
[1]: #Use the following code to install any required libraries that are missing
#!pip install -U pip setuptools wheel
#!pip install -U spacy
#!python -m spacy download en_core_web_sm
#!pip install --upgrade gensim
#!pip install --user torch torchvision torchaudio
#!pip install --user transformer
#!pip install --user pyldavis
#!pip install --user wordcloud
#!pip install --user -U nltk
from pprint import pprint
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import torch
from transformers import BertTokenizer
import gensim.models.phrases as gen
from gensim.models.coherencemodel import CoherenceModel
from wordcloud import WordCloud
import seaborn as sns
import pyLDAvis
import pyLDAvis.sklearn
import re
from gensim import models, corpora
import nltk
nltk.download('stopwords')
import string
string.punctuation
stopwords = nltk.corpus.stopwords.words('english')
from sklearn.decomposition import PCA
from sklearn.cluster import MiniBatchKMeans
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.cluster import KMeans
from sklearn import metrics
from scipy.spatial.distance import cdist

from sklearn.manifold import TSNE from sklearn.decomposition
import LatentDirichletAllocation from nltk.tokenize import
word_tokenize from nltk.stem.snowball import SnowballStemmer from
sklearn.feature_extraction.text import TfidfVectorizer,
```

```

CountVectorizer from sklearn.feature_extraction.text import
CountVectorizer, TfidfVectorizer from tqdm._tqdm_notebook import
tqdm_notebook, tnrange, tqdm from collections import
Counter, OrderedDict from nltk.stem.wordnet import
WordNetLemmatizer lmtzr = WordNetLemmatizer() import warnings
import pyLDAvis.gensim_models import re, nltk, spacy, gensim from
sklearn.decomposition import LatentDirichletAllocation,
TruncatedSVD from sklearn.model_selection import GridSearchCV
import warnings warnings.filterwarnings('ignore')

```

C:\Users\abuwa\anaconda3\lib\sitepackages\sklearn\linear_model_least_angle.py:34: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use `float` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here. Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
method='lar', copy_X=True, eps=np.finfo(np.float).eps,
```

C:\Users\abuwa\anaconda3\lib\sitepackages\sklearn\linear_model_least_angle.py:164: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use `float` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here. Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
method='lar', copy_X=True, eps=np.finfo(np.float).eps,
```

C:\Users\abuwa\anaconda3\lib\sitepackages\sklearn\linear_model_least_angle.py:281: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use `float` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here. Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
eps=np.finfo(np.float).eps, copy_Gram=True, verbose=0,
```

C:\Users\abuwa\anaconda3\lib\sitepackages\sklearn\linear_model_least_angle.py:865: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use `float` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here. Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
eps=np.finfo(np.float).eps, copy_X=True, fit_path=True,
```

C:\Users\abuwa\anaconda3\lib\sitepackages\sklearn\linear_model_least_angle.py:1121: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use `float` by

itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here. Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```

    eps=np.finfo(np.float).eps, copy_X=True, fit_path=True,
C:\Users\abuwa\anaconda3\lib\sitepackages\sklearn\linear_model\_least
_angle.py:1149: DeprecationWarning: `np.float` is a deprecated alias
for the builtin `float`. To silence this warning, use `float` by
itself. Doing this will not modify any behavior and is safe. If you
specifically wanted the numpy scalar type, use `np.float64` here.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
    eps=np.finfo(np.float).eps, positive=False):
C:\Users\abuwa\anaconda3\lib\sitepackages\sklearn\linear_model\_least
_angle.py:1379: DeprecationWarning: `np.float` is a deprecated alias
for the builtin `float`. To silence this warning, use `float` by
itself. Doing this will not modify any behavior and is safe. If you
specifically wanted the numpy scalar type, use `np.float64` here.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
    max_n_alphas=1000, n_jobs=None, eps=np.finfo(np.float).eps,
C:\Users\abuwa\anaconda3\lib\sitepackages\sklearn\linear_model\_least
_angle.py:1621: DeprecationWarning: `np.float` is a deprecated alias
for the builtin `float`. To silence this warning, use `float` by
itself. Doing this will not modify any behavior and is safe. If you
specifically wanted the numpy scalar type, use `np.float64` here.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
    max_n_alphas=1000, n_jobs=None, eps=np.finfo(np.float).eps,
C:\Users\abuwa\anaconda3\lib\sitepackages\sklearn\linear_model\_least
_angle.py:1755: DeprecationWarning: `np.float` is a deprecated alias
for the builtin `float`. To silence this warning, use `float` by
itself. Doing this will not modify any behavior and is safe. If you
specifically wanted the numpy scalar type, use `np.float64` here.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
    eps=np.finfo(np.float).eps, copy_X=True, positive=False):
C:\Users\abuwa\anaconda3\lib\site-
packages\sklearn\decomposition\_lda.py:28: DeprecationWarning:
`np.float` is a deprecated alias for the builtin `float`. To silence
this warning, use `float` by itself. Doing this will not modify any
behavior and is safe. If you specifically wanted the numpy scalar
type, use `np.float64` here.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations EPS
= np.finfo(np.float).eps

```

```
C:\Users\abuwa\anaconda3\lib\sitepackages\sklearn\feature_extraction\
image.py:172: DeprecationWarning: `np.int` is a deprecated alias for
the builtin `int`. To silence this warning, use `int` by itself.
Doing this will not modify any behavior and is safe. When replacing
`np.int`, you may wish to use e.g. `np.int64` or `np.int32` to
specify the precision. If you wish to review your current use, check
the release note link for additional information.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
dtype=np.int):
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\abuwa\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
<ipython-input-1-06ecb96179f3>:43: TqdmDeprecationWarning: This
function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.*` instead of
`tqdm._tqdm_notebook.*` from tqdm._tqdm_notebook import
tqdm_notebook,tnrange,tqdm
```

0.1 Task 1 - What is the distribution of journals per source?

Due to large dataset, limited memory and performance availability, a portion of the dataset it used (220,000 rows). The 'nrows' value can be easily changed for a larger dataset if more memory and performance is available.

```
[2]: df = pd.read_csv("metadata.csv",nrows=220000)
```

Getting the journal count by source and visualizing the data via table.

```
[3]: new_df = df.value_counts(subset=['source_x', 'journal']).
      ↪ reset_index(name="count")
      new_df.head(20)
```

```
[3]:
```

	source_x	journal	count
0	Medline	BMJ	2862
1	Medline	Journal of virology	2154
2	BioRxiv	bioRxiv	2036
3	PMC	Reactions Weekly	1602
4	Medline	Nature	1301
5	Medline	JAMA	1041
6	PMC	PLoS One	1004
7	Medline	Surgical endoscopy	974
8	Medline	Science	789
9	Medline	The New England journal of medicine	766
10	Medline	Proceedings of the National Academy of Science...	665
11	Medline	The Journal of general virology	607
12	PMC	Sci Rep	572

13 Medline Journal of neurointerventional surgery 536

14 Medline AJNR. American journal of neuroradiology 532
15 Medline The Veterinary record 527
16 Medline BMJ open 512
17 Medline PloS one 504
18 Medline Advances in experimental medicine and biology 504
19 PMC Computational Science and Its Applications - I...500

Getting the top 10 most frequent journals that appear in the dataset.

```
[4]: df['journal'].value_counts()[:10]
```

```
[4]: BMJ 3155  
Journal of virology 2154  
bioRxiv 2043  
Reactions Weekly 1620  
Nature 1608  
PLoS One 1276  
JAMA 1134  
Surgical endoscopy 974  
Science 878  
Lancet 875  
Name: journal, dtype: int64
```

Extracting the top 10 most frequent journals from the dataset

```
[5]: top_ten_journals = ['BMJ', 'Journal of virology', 'bioRxiv', 'Reactions_Weekly', 'Nature', 'PLoS One', 'JAMA', 'Surgical endoscopy', 'Science', 'Lancet']
```


Visualizing the dataset via a table after extracting top 10 most frequent journals

```
[6]: new_df = new_df[new_df.journal.isin(top_ten_journals)]
```

Distribution of journals by source via clustered column chart. Plotting is conducted using seaborn catplot.

```
[7]: # Visualization of a barplot
g = sns.catplot(x="source_x", =new_df['count'], hue="journal",
                data=new_df, kind="bar",
                height=5, palette="muted", legend=True, log=True)

g.fig.set_figwidth(18)
g.fig.set_figheight(8)
plt.title('Distribution of publications in 10 most frequent journals ')
plt.xlabel('Source')
plt.ylabel('Distribution')
plt.show()
```



Outcome: Although 7 sources exist in total, the clustered column chart shows only 4 sources. This may be because the remaining sources do not have any journal distribution or they do not possess enough journals that could be appropriately visualized in the chart.

0.1.1 Task 2 - What are the main clusters?

```
[8]: df = pd.read_csv("metadata.csv", nrows=1500)
```

Due to limited memory and performance availability, a portion of the dataset will be used (1500 rows). The 'nrows' value can be easily changed for a larger dataset if more memory and performance is available. A larger dataset will always yield a more accurate representation of the entire dataset.

Dropping all empty rows in the abstract column and extract the column for further analysis.

```
[9]: df = df.dropna(subset=['abstract'])
df['abstract'] = df['abstract'].astype(str)
df = pd.DataFrame(df['abstract'])
df
```

```
[9]:                                     abstract
0      OBJECTIVE: This retrospective chart review des...
1      Inflammatory diseases of the respiratory tract...
2      Surfactant protein-D (SP-D) participates in th...
3      Endothelin-1 (ET-1) is a 21 amino acid peptide...
4      Respiratory syncytial virus (RSV) and pneumoni...
...
1494 Glycyrrhizic acid (GA) is a triterpene glycosi...
1495 Background: Upper respiratory tract infections...
1496 BACKGROUND: Various pathways have been implica...
1497 INTRODUCTION: There is a hyperoxidative state ...
1498 INTRODUCTION: The aim of this study was to inv...

[1435 rows x 1 columns]
```

Verifying that no null values exist in the abstract column.

```
[10]: df.isnull().sum()
```

```
[10]: abstract    0
      dtype: int64
```

Text preprocessing is conducted. We'll also add our own custom stopwords which include some common words used in research papers.

```
[11]: custom_stop_words = [
    'reserved', 'peer', 'CZI', 'reviewed', 'org', 'et', 'author',
    'figure',
    'rights', 'permission', 'Elsevier', 'biorxiv', 'https', 'copyright',
    'medrxiv', 'license', 'preprint', 'fig', 'fig.',
    'al.', 'al', 'PMC', 'doi'
]

stopwords.extend(custom_stop_words)

# A function to prepare the
text def clean(text):
    regex = re.compile('[' + re.escape(string.punctuation) + '0-9\\r\\t\\n]')
    text = regex.sub(" ", text.lower())
    words = text.split(" ")
    words = [re.sub('\\S*@\S*\s?', '', sent) for
```

```

sent in words] words = [re.sub('\s+', ' ', sent) for sent in
words] words = [re.sub("\'", "'", sent) for sent in words] words
= [w for w in words if not len(w) < 2] words = [w for w in
words if w not in stopwords] words = [lmtzr.lemmatize(w) for w
in words] words = ' '.join([str(w) for w in words]) return
words

```

Applying above function to the abstract column.

```
[12]: df["clean_abstract"] = df["abstract"].apply(clean)
```

Tokenizing the text in abstract column with BertTokenizer [13]:

```

def sen_to_vec(sentence):
tokenizer=BertTokenizer.from_pretrained('bert-base-uncased')
tokens=tokenizer.tokenize(sentence) tokens = ['[CLS]'] +
tokens + ['[SEP]']
T=624
padded_tokens=tokens + ['[PAD]' for _ in range(T-len(tokens))]
attn_mask=[ 1 if token != '[PAD]' else 0 for token in
padded_tokens ]
seg_ids=[0 for _ in range(len(padded_tokens))]
sent_ids=tokenizer.convert_tokens_to_ids(padded_tokes
ns) return np.array(sent_ids)

```

After tokenization is completed, it is stored in a new data frame 'final_df' for further analysis.

```

[14]: df["array"] =
df["clean_abstract"].apply(sen_to_vec)
final_df = pd.DataFrame(df["array"]) final_df
= final_df.pop('array').apply(pd.Series)
final_df = final_df.fillna(0)
final_df.head(3)

```

```

[14]: 0      1      2      3      4      5      6      7      8      9      ... \
0  101 7863 15354 3673  3319  5577  4958  5178  4328  6779 ...
1  101 201874295 16464 12859 4141  3378  8319  2537  9152 ...
2  101 14175 18908      4630  5250 11867 17257 25605 3433 15938 ...

      614 615 616 617 618 619 620 621 622 623
0      0   0   0   0   0   0   0   0   0   0
1      0   0   0   0   0   0   0   0   0   0
2      0   0   0   0   0   0   0   0   0   0
[3 rows x 624 columns]

```

Using PCA to reduce the dimensionality of tokenized array data to 2 so that it can be visualized efficiently.


```
[15]: from sklearn.decomposition import PCA
```

```
pca = PCA(n_components = 2)
X_PCA=pca.fit_transform(final_df)

print(final_df.shape)
print(X_PCA.shape)
```

```
(1435, 624)
```


```
(1435, 2)
```

Finding the best K value for our Kmeans model.

```
[16]: distortions = []
K = range(1,30)
for k in K:
    kmeanModel = KMeans(n_clusters=k)
    kmeanModel.fit(X_PCA)
    distortions.append(kmeanModel.inertia_)
```

Finding the best K value by using Elbow method visualization.

```
[17]: plt.figure(figsize=(16,8))
plt.plot(K, distortions, 'bx-')
plt.xlabel('k')
plt.ylabel('Distortion')
plt.title('The Elbow Method showing the optimal k ')
plt.show()
```




By analyzing above graph, we can deduce $k = 3$.

```
[18]: k = 3
kmeans = KMeans(n_clusters=k, random_state=42)
y_pred = kmeans.fit_predict(X_PCA)
df['Clusters'] = y_pred
```

Vizualisation of clusters

```
[19]: sns.countplot(df['Clusters'])
```

```
[19]: <AxesSubplot:xlabel='Clusters', ylabel='count'>
```




2D representation of each document in its corresponding clusters. The clustered are represented by its following colours.

```
[20]: # sns settings
sns.set(rc={'figure.figsize': (13,9)})

# colors
palette = sns.hls_palette(3,  .4,  .9)

# plot
sns.scatterplot(X_PCA[:,0], X_PCA[:,1], hue=y_pred, legend='full', 
               palette=palette)
plt.title('PCA with Kmeans Labels')
plt.savefig("clusterRevised.png")
plt.show()
```



0.1.2 Task 3 - For each cluster, what are the most representative words?

```
[21]: wordcloud = _
      → WordCloud(background_color='white', stopwords=stopwords, max_words=200, max_font_size=40, random_state=42)
      → generate(str(df['clean_abstract']))
```

```
[22]: print(wordcloud)
      fig = plt.figure(1)
      plt.imshow(wordcloud)
      plt.axis('off')
      plt.show()
      fig.savefig("word1.png", dpi=900)
```

We will use cleaned abstract data obtained earlier to produce the wordcloud.

<wordcloud.wordcloud.WordCloud object at 0x000001905114BAF0>



Note: Upon analyzing the above Wordcloud, we can see that even more meaningful results can be obtained if more custom stopwords are added to our list such as “introduction”, “background”, “objective”. This can easily be done by adding these words to our already created list above “custom_stop_words”. However, as the program takes many hours to run for even small datasets, it was decided to not perform this operation again.

0.1.3 Task 4 - What are the most common topics?

Preprocessing of data using gensim.

```
[23]: data =
      df.clean_abstract.values.tolist() def
      sent_to_words(sentences):
          for sentence in sentences:
              yield(gensim.utils.simple_preprocess(str(sentence),
                  deacc=True) #
          ,deacc=True removes punctuations
      data_words =
      list(sent_to_words(data))
      print(data_words[:1])

[['objective', 'retrospective', 'chart', 'review', 'describes',
'epidemiology',
'clinical', 'feature', 'patient', 'culture', 'proven', 'mycoplasma',
'pneumoniae', 'infection', 'king', 'abdulaziz', 'university',
'hospital',
'jeddah', 'saudi', 'arabia', 'method', 'patient', 'positive',
'pneumoniae',
```

```

'culture', 'respiratory', 'specimen', 'january', 'december',
'identified',
'microbiology', 'record', 'chart', 'patient', 'result', 'patient',
'identified',
'required', 'admission', 'infection', 'community', 'acquired',
'infection',
'affected', 'age', 'group', 'common', 'infant', 'pre', 'school',
'child',
'occurred', 'year', 'round', 'common', 'fall', 'spring', 'three',
'quarter',
'patient', 'comorbidities', 'twenty', 'four', 'isolates',
'associated',
'pneumonia', 'upper', 'respiratory', 'tract', 'infection',
'bronchiolitis',
'cough', 'fever', 'malaise', 'common', 'symptom', 'crepitation',
'wheeze',
'common', 'sign', 'patient', 'pneumonia', 'crepitation', 'bronchial',
'breathing', 'patient', 'likely', 'non', 'patient', 'present',
'pneumonia',
'versus', 'patient', 'pneumonia', 'uneventful', 'recovery',
'recovered',
'following', 'complication', 'died', 'pneumoniae', 'infection',
'died', 'due',
'underlying', 'comorbidities', 'patient', 'died', 'pneumoniae',
'pneumonia',
'comorbidities', 'conclusion', 'result', 'similar', 'published',
'data',
'except', 'finding', 'infection', 'common', 'infant', 'preschool',
'child',
'mortality', 'rate', 'pneumonia', 'patient', 'comorbidities',
'high']]

```

Lemmatization

```

[24]: def lemmatization(texts, allowed_postags=['NOUN', 'ADJ', 'VERB',
'ADV']):
    ↪# 'NOUN', 'ADJ', 'VERB',
    ↪'ADV' texts_out = []
    for sent in texts:
        doc = nlp(" ".join(sent))
        texts_out.append(" ".join([token.lemma_ if token.lemma_ not
in_
    ↪['-PRON-'] else '' for token in doc if token.pos_ in
allowed_postags])) return texts_out

```

```
[25]: # Initialize spacy 'en' model, keeping only tagger component (for efficiency)
```

```
# Run in terminal: python -m spacy download en
nlp = spacy.load("en_core_web_sm", disable=['parser', 'ner']) #
Do lemmatization keeping only Noun, Adj, Verb, Adverb
data_lemmatized = lemmatization(data_words,
allowed_postags=['NOUN', 'ADJ', '_',
                'VERB', 'ADV']) #select noun and verb

print(data_lemmatized[:2])
```

```
['objective retrospective chart review describe epidemiology clinical
feature patient culture prove mycoplasma infection king method
patient positive pneumoniae culture respiratory specimen identify
microbiology record chart patient result patient identify require
admission infection community acquire infection affect age group
common infant pre school child occur year round common fall spring
quarter patient comorbiditie isolate associate upper respiratory
tract cough fever malaise common symptom crepitation wheeze common
sign patient pneumonia bronchial breathing patient likely non patient
present pneumonia patient pneumonia uneventful recovery recover
follow complication die pneumoniae infection die due underlie
comorbiditie patient die pneumonia comorbiditie conclusion result
similar publish datum find infection common infant preschool child
mortality rate pneumonia patient comorbiditie high', 'inflammatory
disease respiratory tract commonly associate elevated production
nitric oxide increase index dependent oxidative stress know anti
microbial anti inflammatory anti oxidant property various line
evidence support contribution lung injury several disease model basis
biochemical evidence often presume dependent oxidation due formation
oxidant peroxyxynitrite alternative mechanism involve phagocyte derive
heme protein myeloperoxidase peroxidase operative condition
inflammation overwhelm literature generation activity respiratory
tract scope commentary review area comprehensively instead focus
recent evidence concept presume contribution inflammatory disease
lung']
```

Data is vectorized using CountVectorizer

```
[26]: vectorizer = CountVectorizer(analyzer='word',
                                min_df=10,
                                # minimum reqd occurrences of a word
                                stop_words='english',
                                # remove stop words
                                lowercase=True,
                                # convert all words to lowercase
                                token_pattern='[a-zA-Z0-9]{3,}')

data_vectorized = vectorizer.fit_transform(data_lemmatized)
```

Building an LDA Model.

Since we had 3 clusters we will stick to 3 topics, to be able to make a comparison.

```
[27]: # Build LDA Model
lda_model = LatentDirichletAllocation(n_components=3,          # Number of _
    → topics
                                max_iter=10,
    # Max learning iterations
                                learning_method='online',
                                random_state=100,
    # Random state
                                batch_size=128,
    # n docs in each learning iter
                                evaluate_every = -1,
    # compute perplexity every n iters, default: Don't
                                n_jobs = -1,
    # Use all available CPUs
                                )
lda_output = lda_model.fit_transform(data_vectorized)
print(lda_model)  # Model attributes
```

```
LatentDirichletAllocation(learning_method='online', n_components=3,
                          n_jobs=-1, random_state=100)
```

For each topic, get the top n keywords.

```
[29]: # Show top n keywords for each topic def
show_topics(vectorizer=vectorizer, lda_model=lda_model,
n_words=20):
    keywords = np.array(vectorizer.get_feature_names())
    topic_keywords = [] for topic_weights
    in lda_model.components_:
        top_keyword_locs = (-
        topic_weights).argsort()[::-n_words]
        topic_keywords.append(keywords.take(top_keyword_locs))
    return topic_keywords
```

```

topic_keywords = show_topics(vectorizer=vectorizer,
                              lda_model=lda_model, n_words=15)

# Topic - Keywords Dataframe df_topic_keywords =
pd.DataFrame(topic_keywords) df_topic_keywords.columns = ['Word
'+str(i) for i in range(df_topic_keywords.
→shape[1])] df_topic_keywords.index = ['Topic '+str(i) for i in
range(df_topic_keywords.
→shape[0])]
df_topic_keywords

```

```

[29]:      Word 0      Word 1      Word 2      Word 3      Word 4      Word 5 \
Topic 0      use    disease influenza    health      virus sequence
Topic 1 patient infection    study influenza respiratory clinical
Topic 2    cell    protein    virus    gene    infection    viral

      Word 6 Word 7 Word 8 Word 9 Word 10 Word 11 Word 12 \
Topic 0    method study    result model    datum pandemic analysis
Topic 1      high    day    severe result treatment disease mortality
Topic 2 expression    use activity human    study    induce response

      Word 13      Word 14
Topic 0 transmission population
Topic 1      virus      case
Topic 2      result      host

```

Adding appropriate labels for the 3 topics.

```

[30]: Topics = ["Research", "Disease and Treatment", "Virus Behaviour"]

df_topic_keywords["Topics"]=Topics
df_topic_keywords

```

```

[30]:      Word 0      Word 1      Word 2      Word 3      Word 4      Word 5 \
Topic 0      use    disease influenza    health      virus sequence
Topic 1 patient infection    study influenza respiratory clinical
Topic 2    cell    protein    virus    gene    infection    viral

      Word 6 Word 7 Word 8 Word 9 Word 10 Word 11 Word 12 \
Topic 0    method study    result model    datum pandemic analysis
Topic 1      high    day    severe result treatment disease mortality
Topic 2 expression    use activity human    study    induce response

      Word 13      Word 14      Topics
Topic 0 transmission population    Research Topic 1
      virus case Disease and Treatment

```


Topic 2 result host Virus Behaviour

Defining a function to predict topic for a given text document. Furthermore, as an example, topic is inferred for a given text. This can be seen in the result below.

```
[36]: nlp = spacy.load('en_core_web_sm', disable=['parser',
'ner']) def predict_topic(text, nlp=nlp):
    global sent_to_words
    global lemmatization
    # Step 1: Clean with simple preprocess
    mytext_2 = list(sent_to_words(text))
    # Step 2: Lemmatize mytext_3 = lemmatization(mytext_2,
    allowed_postags=['NOUN', 'ADJ', 'VERB', '_', 'ADV'])
    # Step 3: Vectorize transform mytext_4
    = vectorizer.transform(mytext_3)
    # Step 4: LDA Transform topic_probability_scores =
    lda_model.transform(mytext_4) topic =
    df_topic_keywords.iloc[np.argmax(topic_probability_scores), 1:14].
    .values.tolist()

    # Step 5: Infer Topic
    infer_topic =
    df_topic_keywords.iloc[np.argmax(topic_probability_scores), _
    ,-1]

    #topic_guess =
    df_topic_keywords.iloc[np.argmax(topic_probability_scores), _
    ,Topics] return infer_topic, topic,
    topic_probability_scores
    # Predict the topic mytext = ["initial coronavirus study
    promising we will now finish final_
    ,analysis"] infer_topic, topic, prob_scores =
    predict_topic(text = mytext) print(infer_topic)
```

Research

0.1.4 Task 5 - What are the most common topics in each cluster?

Information from the topics in task 4 is used to characterise the topics in each cluster and adding it to the dataframe.

```
[37]: def apply_predict_topic(text):
    text = [text]
    infer_topic, topic, prob_scores = predict_topic(text = text)
    return(infer_topic)
df["Topic_key_word"] = df['abstract'].apply(apply_predict_topic)
df.head()
```

```

[37]:                                     abstract \
0 OBJECTIVE: This retrospective chart review des...
1 Inflammatory diseases of the respiratory tract...
2 Surfactant protein-D (SP-D) participates in th...
3 Endothelin-1 (ET-1) is a 21 amino acid peptide...
4 Respiratory syncytial virus (RSV) and pneumoni...

                                     clean_abstract \
0 objective retrospective chart review describes...
1 inflammatory disease respiratory tract commonl...
2 surfactant protein sp participates innate resp...
3 endothelin amino acid peptide diverse biologic...
4 respiratory syncytial virus rsv pneumonia viru...

                                     array Clusters \
0 [101, 7863, 15354, 3673, 3319, 5577, 4958, 517... 0
1 [101, 20187, 4295, 16464, 12859, 4141, 3378, 8... 2
2 [101, 14175, 18908, 4630, 5250, 11867, 17257, ... 1
3 [101, 2203, 14573, 18809, 13096, 5648, 25117, ... 2
4 [101, 16464, 26351, 22123, 4818, 7865, 12667, ... 2

      Topic_key_word
0      Disease and Treatment
1      Virus Behaviour
2      Virus Behaviour
3      Virus Behaviour
4      Virus Behaviour

```

Visualization of top 3 common topics.

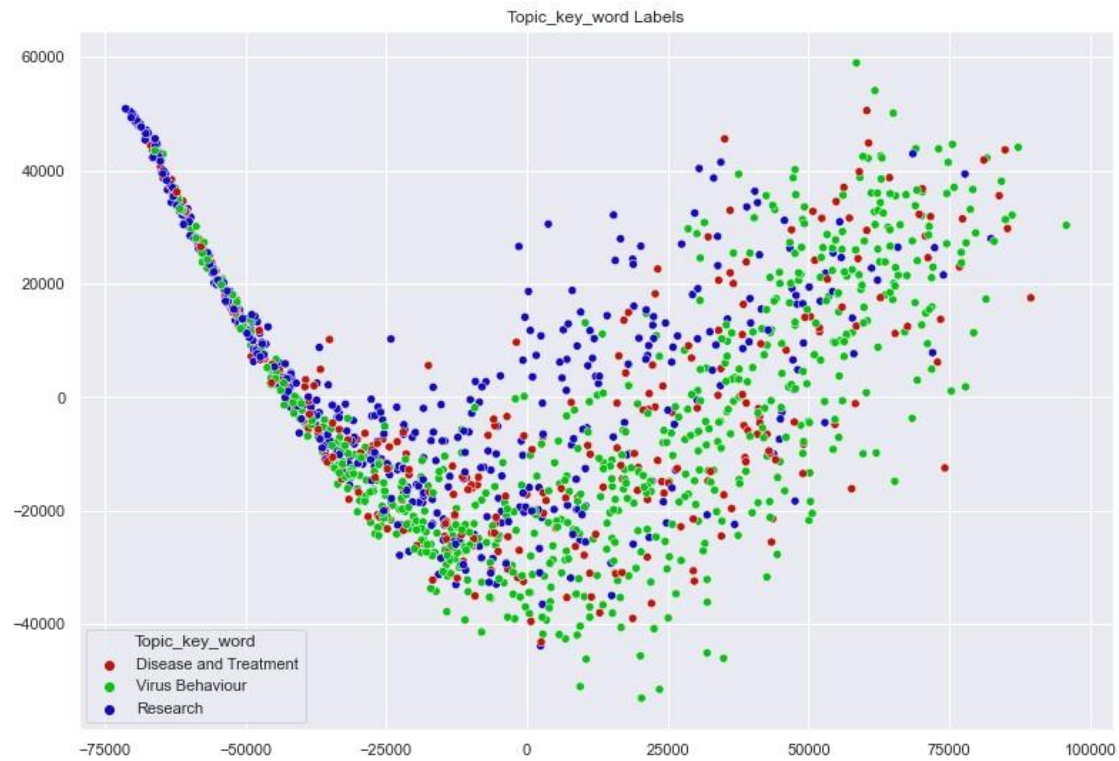
```

[38]: # sns settings
sns.set(rc={'figure.figsize':(14,10)})

# colors
palette = sns.hls_palette(3,  =.4,  =.9)

# plot
sns.scatterplot(X_PCA[:,0], X_PCA[:,1], hue=df['Topic_key_word'], _
→ legend='full', palette=palette)
plt.title('Topic_key_word Labels')
plt.savefig("clusterRevised.png")
plt.show()

```



Comparison of Kmeans clusters and Topics from Lda

```
[39]: fig, ax =plt.subplots(1,2)
sns.countplot(df['Topic_key_word'], ax=ax[0])
sns.countplot(df['Clusters'], ax=ax[1])
fig.show()
```

